

Mockup Test — Research Synthesis

n=4 (Banks, Natali, Nicole, Priya) | Trust hypotheses T1–T5 | May 2026

Five trust questions

These are not generic usability findings. Each one tests a specific failure mode identified before testing began.

T1 — Does the AI-generated scope feel credible enough to act on?

Yes. All four approved without editing. Credibility came from structure and the local benchmark, not the 92% confidence score — which was missed by Banks and questioned by Nicole and Priya ("what is this based on?"). The score is decoration right now.

Verdict:

T1 passes, conditionally. Credibility comes from structure and comprehensiveness. The confidence score isn't carrying its weight.

T2 — Does "verified" translate into felt safety?

Partially. The flagged insurance on Quickfix landed as a real warning for 3/4 — all three dropped that contractor. License and permit counts didn't shift confidence on their own. Users understand what verification means in real life; they don't know what "47 permits pulled" tells them about their repair. Needs plain-language translation, not raw data.

Verdict:

T2 partially passes. Warning signals work when seen. Positive verification signals need translation into plain risk language to do real work.

T3 — Does the recommendation feel like advice or a list?

Depends on whether they saw the reasoning. "Why we picked these three" worked — Priya, Nicole, and Natali read it and felt more confident. Banks never saw it and assumed Jason "paid the backend to be promoted." The logic is there. The hierarchy buries it.

Verdict:

T3 passes when the reasoning is visible, fails when it isn't. The block needs to lead the section, not close it.

T4 — Will users approve outreach before seeing proof it works?

Mostly yes. Banks, Natali, Nicole approved straight through. Priya hesitated and clicked Edit scope first — she didn't want contractors receiving AI-generated content before she reviewed it. Three of four approved with no friction. One surfaced the exact failure mode that matters most.

Verdict:

T4 mostly passes, but Priya's hesitation is the canary. The CTA needs to feel like a deliberate review moment, not just a proceed button.

T5 — Is three quotes enough?

Enough to decide, not enough to commit. All four picked a contractor. None asked for more quotes. But Nicole, Natali, and Priya all said they'd want to contact the contractor before booking — to confirm cost, materials, access. Three quotes solves comparison. It doesn't solve the final-mile anxiety of hiring someone you've never spoken to.

Verdict:

T5 passes on quantity. Fails on what follows selection. The product treats booking as the endpoint; users treat it as the beginning of a relationship they haven't started.

Top 3 friction points

1. No way to contact the contractor before booking — 4/4 sessions

Every participant wanted a message field, contact details, or direct access before committing. Natali wanted contact details and a note field. Nicole's single requested change was "add an easy way to message or call the contractor before final booking." Priya said she'd call first. Banks is the only one who'd act without it — and he's also the user who assumed Jason paid to be promoted.

2. "Why we picked these three" is invisible — missed by 2/4 unprompted

Banks missed it entirely and suggested repositioning. The other three found it only after reading past the table. It's doing the most trust work on the screen and it's buried below where decisions are already forming.

3. Confidence score and ranking logic unexplained — 3/4 sessions

Nicole and Priya both asked what 92% is based on. Banks didn't notice it at all. Priya's single requested change was "add clearer explanations for how confidence scores and contractor

rankings are calculated." A number without a source reads as reassurance theater.

Per-screen findings

Onboarding

No abandonment. All four said "about right" on length. Nicole and Priya both hesitated on address and year built — not sure whether the data improved recommendations. Add one-line explanations per field to remove this friction.

Reactive intake

No one balked at three photos. Natali found photo-driven intake actively reassuring — the AI analyzing photos reduced her need to diagnose the problem herself. Priya appreciated guidance on what each photo should show. Nicole and Priya both wanted the ability to correct AI auto-labels if wrong. The feature feels trustworthy; users want an edit escape hatch.

Scope screen

Strongest screen in the product. All four approved. The local benchmark (\$180–\$320) was the most cited trust element across all sessions. Itemized labor and acceptance criteria felt professional and specific. Unit-priced add-ons: zero participants mentioned them. Either invisible or not processed — needs a direct probe in Week 8.

Contractor matching

Eyes go to AI TOP MATCH first, then rating and insurance. Insurance flag noticed by 3/4; missed by Banks. "Why we picked these three" read by 3/4; missed by Banks. Permit counts and license numbers read as data, not protection. Priya: "reviews alone are not enough, but I need these signals translated, not just shown."

Quote comparison

Plain-language summary read and valued by all four. OUTLIER/BETTER tags noticed by Nicole and Priya. Quickfix being incomplete made all four uncomfortable — read as a proxy for contractor reliability, not just missing info. Jason was the unanimous pick across all sessions.

Screen works. The gap is what happens after the pick.

Dashboard

Positive reaction across all four. But every participant immediately qualified their comfort: "as long as I can step in," "as long as contact details are accessible," "as long as I can override." The autonomy thesis landed. The condition attached to it is universal.

Top 5 quotes

Banks, contractor matching

"Jason seems like he paid the backend to be promoted."

Priya, scope screen

"The price range helps because otherwise I have no idea what normal is."

Natali, wrap-up

"A final decision stays after me."

Nicole, quote comparison

"Quickfix is cheaper, but the missing details make me nervous."

Priya, contractor matching

"AI top match is useful, but I need to know why it thinks that."

Recommended changes

1. Add a contractor contact or note field before booking

Message field or visible contact info. Requested by 4/4 participants. Highest-stakes moment in the flow.

2. Move "Why we picked these three" above the comparison table

Users form opinions while scanning the table. The reasoning needs to come first, not after.

3. Explain confidence score

One line: what it's based on. Three sessions raised this. A number without a source doesn't build trust.

4. Make the flagged insurance impossible to miss

Amber text in a table row isn't enough. Instead, a warning banner above the contractor card. A user who books a contractor with lapsed insurance because the flag wasn't visible is a liability miss.

5. Make "Approve & find contractors" feel like a review moment, not a proceed button

Add copy that acknowledges outreach is about to be sent. Priya's hesitation here was the clearest signal this CTA needs weight.

Open questions for Week 8

1. What happens when the AI is wrong?

Every session ran on a clean scope. The silent failure mode — AI is confidently wrong, user doesn't catch it — was never tested. Week 8 needs at least one scope with a plausible error.

2. Do users understand unit-priced add-ons?

Zero participants mentioned them. Core to the apples-to-apples promise — needs a direct probe.

3. What happens during the wait?

Users assumed contractors would respond. What does trust look like if a contractor doesn't respond within 24 hours?

4. Does the annual pricing model work for low-frequency users?

Banks wouldn't pay \$149/year. Natali and Priya preferred monthly. The frequency assumption under the subscription model needs a pricing test.

5. What do first-time homeowners do differently?

All four participants had prior contractor experience. The no-reference-point user was not covered.

AI alone, or human in the loop?

Not alone. Not a credentialed expert. A reachable contractor.

All four accepted AI judgment for organizing, diagnosing, comparing, and recommending. None wanted to do those things themselves. But every participant drew a line before full handoff — at the contractor's contact details, at the message field, at the ability to review before outreach fired.

"A final decision stays after me." — Natali

Users trust AI to reduce the decision surface. They want human access at the moment of commitment. The gap in the current mockup is that booking happens without a single word exchanged with the contractor. That's where trust breaks down for 3/4 participants.

The fix isn't adding a human to the Homewise team. It's making the contractor feel like a person the user is about to hire — not a slot the AI filled.